

# Introduction to R

---

Instructor: Christian Capezza

Course leader: Prof. Biagio Palumbo

Course: Statistics for Technology, a.y. 2019/2020

MSc in Mechanical Engineering for Design and Production

17 October 2019

[christian.capezza@unina.it](mailto:christian.capezza@unina.it)

Introduction

The very basics

Data structures

Access to elements in the data

Work with data: commonly used functions and operations

Loops

Work with data stored in a file

Plots: the very basics

# Introduction

---

# Why use R?

- Programming language for statistics
- Free and open source
- Packages
- Great support (Stackoverflow, R-Bloggers, etc.)
- Many free resources to learn

# Install R and RStudio

Install R before RStudio

- R: <https://www.r-project.org/>
- RStudio: <https://www.rstudio.com/>

Rstudio is an integrated development environment (IDE) for R

First thing to do: open a new script

- File -> New file -> R script
- CTRL/CMD + SHIFT + N

Four main panels

- Upper left: Script
- Lower left: Console
- Upper right: Environment variables, history
- Lower right: Files, plots, help

## Set a new project

When you work with data stored in a file, you want to be able to load the data from the script.

- Create a new folder that will contain your project (scripts, data, plots you will generate, etc.)
- Save the new script into this folder
- Go to Session -> Set Working Directory -> to source file location
  - Alternatively,

```
setwd("C:/New project")
```

# The very basics

---



## Start programming: basic operations

You can write code lines in the console, better in the script.

Note: results appear in the console panel

```
1 + 2 # Sum. Note: everything after # is a comment
```

```
## [1] 3
```

```
5 * 4 # Product
```

```
## [1] 20
```

```
7 / 4 # Division
```

```
## [1] 1.75
```

```
3 ^ 2 # Power
```

```
## [1] 9
```

## Assign values to variables with <-

```
a <- 1.6
```

Note: the variable a appears in the environment panel

Now just type a to show its value

```
a
```

```
## [1] 1.6
```

## Access to variable values and use them

```
a + 3
```

```
## [1] 4.6
```

```
b <- a * 2
```

```
b
```

```
## [1] 3.2
```

```
b <- 5 # You can overwrite b
```

```
b
```

```
## [1] 5
```

```
b <- b + 1 # You can overwrite b using b itself
```

```
b
```

```
## [1] 6
```

# Functions

You can use

- functions available in R
- functions provided by packages
- user-defined functions: you define your own function

A word followed by round brackets calls a function. Example: the function `c()` concatenates elements in a single vector:

```
c(1, 2.5, 4.5)
```

```
## [1] 1.0 2.5 4.5
```

A function is characterized by:

- the name that identifies it (e.g. `c`)
- the required inputs (some are given by default)
- the outputs (in R it is always a single object)

Use `help` to understand how a specific function works. `help` and `?` are equivalent. Example:

```
help(c)
```

```
?c
```

# Data structures

---

# Atomic vectors

Recommended: chapter 'Data structures' from (Wickham 2014).

Four types of atomic vectors:

```
dbl_var <- c(1, 2.5, 4.5) # double
```

```
class(dbl_var)
```

```
## [1] "numeric"
```

```
int_var <- c(1L, 6L, 10L) # integer
```

```
class(int_var)
```

```
## [1] "integer"
```

```
log_var <- c(TRUE, FALSE, T, F) # logical (T/TRUE equivalent)
```

```
class(log_var)
```

```
## [1] "logical"
```

```
chr_var <- c("these are", "some strings") # character
```

```
class(chr_var)
```

```
## [1] "character"
```

# Factors

Categorical variable. Repeated values are considered as the same level. You can't use values that are not in the levels

```
x <- factor(c("a", "b", "a"))  
x  
## [1] a b a  
## Levels: a b  
class(x)  
## [1] "factor"  
levels(x)  
## [1] "a" "b"  
table(x) # absolute frequencies  
## x  
## a b  
## 2 1  
table(x) / length(x) # relative frequencies
```



## Numerical vectors: alternative ways to define them

```
c(6.4, 4.3, 7, 3) # as seen before
```

```
## [1] 6.4 4.3 7.0 3.0
```

```
1:10 # Sequence of numbers increasing by one
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
seq(from = 1, to = 10, by = 3)
```

```
## [1] 1 4 7 10
```

```
seq(from = 2, to = 10, length = 5)
```

```
## [1] 2 4 6 8 10
```

```
rep(3, length = 5)
```

```
## [1] 3 3 3 3 3
```

# Matrices

Different ways to define a matrix

```
M <- matrix(data = c(11, 12, 13, 14, 15, 16),  
            nrow = 2, ncol = 3, byrow = FALSE)
```

```
class(M)
```

```
## [1] "matrix"
```

```
matrix(data = c(11, 12, 13, 14, 15, 16),  
       nrow = 2, ncol = 3, byrow = TRUE)
```

```
##      [,1] [,2] [,3]
```

```
## [1,]   11   12   13
```

```
## [2,]   14   15   16
```

# Matrices

```
rbind(c(11, 13, 15), c(12, 14, 16))    # bind vecs as rows
##           [,1] [,2] [,3]
## [1,]      11   13   15
## [2,]      12   14   16

cbind(c(11, 12), c(13, 14), c(15, 16)) # bind vecs as cols
##           [,1] [,2] [,3]
## [1,]      11   13   15
## [2,]      12   14   16
```

## Heterogeneous data structures: list

List, it is a container and can contain objects of any type and any length. Every object can be identified by a name

```
l <- list(num = 1:10, cha = c("a", "b", "c"),  
          fac = factor(c("c1", "c2")))
```

```
class(l)
```

```
## [1] "list"
```

```
l
```

```
## $num
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
##
```

```
## $cha
```

```
## [1] "a" "b" "c"
```

```
##
```

```
## $fac
```

```
## [1] c1 c2
```

```
## Levels: c1 c2
```

## Heterogeneous data structures: data frame

The most important data structure for storing data in R, they combine the behaviour of matrices and lists

- Like matrices, data frames have rows and columns, however
  - matrices require all elements are of the same type
  - data frames allow objects of different type for each column
- Like lists, they contain objects/columns of different types, however
  - lists allow different lengths for each contained object
  - data frames require the same length for each column

# Heterogeneous data structures: data frame

```
df <- data.frame(nome = c("Gennaro Esposito"),
                 eta = c(24, 21, 32),
                 sesso = factor(c("M", "F"))) # Does not work!

df <- data.frame(name = c("Gennaro Esposito", "Maria Rossi"),
                 age = c(24, 21),
                 sex = factor(c("M", "F")))

class(df)

## [1] "data.frame"

df

##           name age sex
## 1 Gennaro Esposito  24  M
## 2      Maria Rossi  21  F
```

## **Access to elements in the data**

---

## Extract elements from vectors

```
a <- c(5, 23, 9)
a[2] # Indicate the position in the vector

## [1] 23

a[c(1, 3)] # Extract multiple elements

## [1] 5 9

a[c(TRUE, TRUE, FALSE)] # Extract only TRUE

## [1] 5 23

a[c(1, 1, 1, 3, 3, 3)] # Extract elements multiple times

## [1] 5 5 5 9 9 9

a[] # Take all

## [1] 5 23 9

b <- 2
a[b] # You can of course use variables as indexes

## [1] 23
```



## Extract elements from vectors

You can also extract by removing undesired elements with minus sign:

```
a[- 2] # take everything except second
```

```
## [1] 5 9
```

```
a[- c(1, 2)] # take everything except first and second
```

```
## [1] 9
```

## Extract elements from matrices

```
M <- rbind(c(11, 13, 15), c(12, 14, 16))
```

```
M
```

```
##      [,1] [,2] [,3]
```

```
## [1,]   11   13   15
```

```
## [2,]   12   14   16
```

```
M[2, 3] # indicate row and column
```

```
## [1] 16
```

```
M[1, ] # take entire row as vector
```

```
## [1] 11 13 15
```

```
M[, 3] # take entire column as vector
```

```
## [1] 15 16
```

```
M[1, , drop = FALSE] # take entire row as matrix
```

```
##      [,1] [,2] [,3]
```

```
## [1,]   11   13   15
```

## Extract elements from list

(Results not shown in the slide)

```
l <- list(num = 1:10, cha = c("a", "b", "c"))  
l[1]           # extract object by position as list  
l[[1]]         # extract object by position as it is  
l["num"]       # extract object by name as list  
l$num          # as above (advantageous in RStudio)  
l[["num"]]     # extract object by name as it is
```

## Extract elements from data frame

It works in the same way as matrices (results not shown in the slide)

```
df <- data.frame(name = c("Gennaro Esposito", "Maria Rossi"),  
                 age = c(24, 21),  
                 sex = factor(c("M", "F")))
```

```
df[2, 3] # indicate row and column
```

```
df[1, ] # take entire row (remains data frame)
```

```
df[, 3] # take entire column as vector
```

```
df$sex # as above (advantageous in RStudio)
```

```
df[, 3, drop = FALSE] # take entire column as data frame
```

## Extract elements from data frame

It works in the same way as lists, too (results not shown in the slide)

```
df[2]           # extract object by position as data frame
df[[2]]         # extract object by position as it is
df["sex"]       # extract object by name as data frame
df[["sex"]]     # extract object by name as it is
```

## Use logical variables to create conditions...

We can find positions in vectors/matrices/lists that satisfy a condition...

```
a <- c(3, 2, 1, 4)
a > 2    # TRUE where condition satisfied, FALSE otherwise
## [1] TRUE FALSE FALSE TRUE
a >= 2 & a <= 3 # AND condition
## [1] TRUE TRUE FALSE FALSE
a == 3 | a == 4 # OR condition
## [1] TRUE FALSE FALSE TRUE
which(a >= 2)    # Gives positions where condition satisfied
## [1] 1 2 4
```

## ... and use them to extract elements

... and use them to extract the corresponding elements

```
a <- c(3, 2, 1, 4)
```

```
a[a > 2]
```

```
## [1] 3 4
```

```
a[a >= 2 & a <= 3]
```

```
## [1] 3 2
```

```
a[a == 3 | a == 4]
```

```
## [1] 3 4
```

```
a[which(a >= 2)] # It is the same as a[a >= 2]
```

```
## [1] 3 2 4
```

# Missing values

In R, NA indicates missing values

```
a <- c(6.7, NA, 4.2)
```

```
is.na(a) # a == NA does not work
```

```
## [1] FALSE TRUE FALSE
```

```
!is.na(a) # ! gives the opposite condition
```

```
## [1] TRUE FALSE TRUE
```

```
which(!is.na(a)) # gives where you have values
```

```
## [1] 1 3
```



## **Work with data: commonly used functions and operations**

---

## Useful functions for vectors: element-wise operations

The following operations are applied element by element, then they return a vector with the same length of a

```
a <- c(5.3, 6, 1.5, 9)
```

```
b <- c(2, 5.1, 4, 1)
```

```
a + b
```

```
2 * a
```

```
a ^ 2 # inverse is a ^ 0.5 or sqrt(a)
```

```
exp(a) # inverse is log(a)
```

```
sin(a) # analogously use cos tan etc.
```

## Useful functions for vectors

The following return a single value (summary statistics)

```
length(a)
a[length(a)] # gives the last element
sum(a) # sum all elements
mean(a) # mean: equivalent to sum(a) / length(a)
sd(a) # Standard deviation (obtained dividing by (N-1))
sum(a ^ 2) # sum of squares
min(a) # max(a)
which.min(a) # which.max(a)
crossprod(a, b) # inner product between two vectors
crossprod(a, a) # equivalent to sum of squares
```

## Useful functions for logical

When mathematical operations are applied to logical, FALSE are converted to 0 and TRUE to 1. Very useful to calculate proportions:

```
a <- c(T, F, F, T, T, T, F, T)
sum(a) # total number of trues
sum(!a) # total number of falses
mean(a) # proportion of trues
mean(!a) # proportion of falses
```

## Useful functions for matrices

```
M <- rbind(c(11, 13, 15), c(12, 14, 16), c(5, 7, 2))
N <- rbind(c(14, 10, 17), c(4, 3, 9), c(12, 24, 8))
nrow(M) # number of rows
ncol(M) # number of columns
dim(M)  # two-dim. vector with n. of rows and cols
M[nrow(M), ] # last row
M[, ncol(M)] # last column
M %*% N # Matrix product (check dimensions!)
M * N # Product elementwise (check dimensions!)
det(M) # determinant
solve(M) # inverse matrix (if possible)
```

## apply: apply function to each row/column

```
rowMeans(M) # mean of each row
apply(M, 1, function(x) mean(x)) # equivalent
apply(M, 1, mean) # equivalent
colMeans(M) # mean of each column
apply(M, 2, mean) # equivalent
rowSums(M) # sum of each row
apply(M, 1, sum) # equivalent
colSums(M) # sum of each column
apply(M, 2, sum) # equivalent
rowSd <- apply(M, 1, sd)
```

## Useful functions for lists and data frames

```
l <- list(num = 1:10, cha = c("a", "b", "c"),  
          fac = factor(c("c1", "c2")))  
names(l) # list element names: useful in loops or to extract  
length(l) # number of objects in list  
l$new # if you assign something, appends to end of list  
summary(l) # summary statistics, useful with data frames  
str(l)      # returns the structure, useful with data frames
```

# lapply: apply function to each element

Works with vectors, lists, and data frames

```
lapply(l, function(x) length(x)) # returns a list
```

```
## $num
```

```
## [1] 10
```

```
##
```

```
## $cha
```

```
## [1] 3
```

```
##
```

```
## $fac
```

```
## [1] 2
```

```
sapply(l, function(x) length(x)) # returns a vector or matrix, i
```

```
## num cha fac
```

```
## 10 3 2
```



# Loops

---

# For loop

Avoids writing code many times: prevents propagation of errors and changes are simpler. After `in` you put the vector over which you want to iterate (results not shown).

```
print(1); print(2); print(3)
for (ii in 1:3) print(ii) # the same as above
for (ii in c(1, 5, 2)) print(ii)
for (name in c("hi", "hey")) print(name) # not only numbers!
```

## Work with data stored in a file

---

# Packages

```
install.packages("MASS") # install a new package  
library(MASS)           # load package to use its functions/data
```

## Load data

Once you load the library ISLR, the data frame Auto is automatically available in the variable Auto

```
library(ISLR)           # Contains Auto data
data(Auto)              # load the data
summary(Auto)
View(Auto)              # View the data.frame in a new sheet
fix(Auto)               # View the data.frame in a new window
```

## Load data from file

The same data are available at the following link <http://faculty.marshall.usc.edu/gareth-james/ISL/Auto.csv> as indicated in (James et al. 2013). Download the csv file and put the file in the working directory (or, better, in a subfolder of the working directory that we call data/). Then, load it to get a data frame

```
df <- read.csv("data/Auto.csv")
str(df)


## 'data.frame':    397 obs. of  9 variables:
## $ mpg          : num  18 15 18 16 17 15 14 14 14 15 ...
## $ cylinders    : int   8  8  8  8  8  8  8  8  8  8 ...
## $ displacement: num   307 350 318 304 302 429 454 440 455 390
## $ horsepower  : Factor w/ 94 levels "?","100","102",...: 17 3
## $ weight       : int  3504 3693 3436 3433 3449 4341 4354 4312
## $ acceleration: num   12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year         : int   70 70 70 70 70 70 70 70 70 70 ...
## $ origin       : int    1  1  1  1  1  1  1  1  1  1 ...
```

## Plots: the very basics

---

## Graphics: one variable (numeric)

```
plot(df$mpg) # Plots n vector values against 1:n
```





## Graphics: one variable (numeric)


Results not shown: try them!

```
plot(df$mpg, type = "l") # Lines instead of points
plot(df$mpg, col = "red") # Set points/lines colours
plot(df$mpg, pch = 16) # Change point shape
plot(df$mpg, size = 3) # Change point/line size
plot(df$mpg, type = "l", lty = 2) # Change line type
plot(df$mpg, main = "Title") # Set plot title
plot(df$mpg, xlab = "x axis") # Set x axis title
plot(df$mpg, ylab = "y axis") # Set y axis title
plot(df$mpg, xlim = c(10, 50)) # Set limits of x axis
plot(df$mpg, ylim = c(0, 60)) # Set limits of y axis
```

# Graphics: one variable (numeric)

```
hist(df$mpg) # histogram of frequencies
```

Histogram of df\$mpg



## Graphics: one variable (numeric)


`probability = TRUE` is useful in cases you want to plot densities over histogram. Note: If you want to add a curve or points over an existing plot, after creating the first plot (e.g. `plot`, or `hist`), use `lines` or `points`.

Results not shown: try them!

```
hist(df$mpg, breaks = 40) # set number of bins
hist(df$mpg, probability = TRUE) # histogram of densities
lines(density(df$mpg)) # adds kernel density estimate
```

## Graphics: one variable (numeric)


```
boxplot(df$mpg)
```



# Graphics: two variables (numeric)

Scatterplot of two variables


```
plot(df$mpg, df$displacement)
```



## Graphics: two variables (one factor one numeric)

Each box plot is the one of all observations of displacement corresponding to the given level of the factor variable.

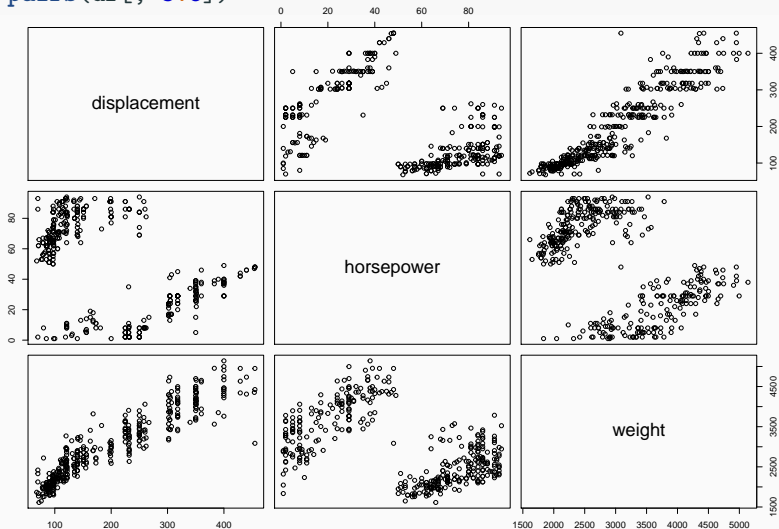
```
plot(factor(df$cylinders), df$displacement)
```



# Graphics: more than two variables (numeric)

Scatterplot of three variables

```
pairs(df[, 3:5])
```



# Random variables

---



## Random variables: normal distribution

`rnorm`, `dnorm`, `pnorm`, `qnorm` are the four functions related to the normal distribution. Analogous functions are available for the other distributions (e.g. `rexp`, `dexp`, `pexp`, `qexp` for the exponential). The four functions, distinguished by the first letter, give

- `dnorm(x)`: give the density  $f_X(x)$  of the random variable  $X$  in  $x$
- `pnorm(x)`: give the cumulative distribution function (cdf)  $F_X(x)$ , i.e. the value of the probability  $Pr(X \leq x)$
- `qnorm(p)`: give the quantile of the random variable  $X$  corresponding to the probability  $p \in [0, 1]$ , i.e. it returns the value  $x$  such that  $F_X(x) = p$ . In other words, it is the inverse of the cdf:  $x = F_X^{-1}(p)$ ,
- `rnorm(n)`: generate  $n$  random observations of  $X$

## Random variables: normal distribution

Be careful that:

- For each distribution you must provide their parameters! For example, for the normal, you must provide mean and standard deviation. Check the help!
- Check also in the help the parametrization of the distribution, sometimes random variables are parametrized in more than one way.


## Plot of a function

Often you are interested in plotting a theoretical parametric density over a histogram, to see if data are fit well. More in general let us see how `plot` works

## Plot a density over a histogram

You can also plot a parametric density over the histogram. Use `curve` with the argument `add = TRUE`.

```
hist(df$mpg, probability = TRUE)  
curve(dnorm(x, mean(df$mpg), sd(df$mpg)), add = TRUE)
```



## How does curve work?

curve allows to build a plot of a function in a compact way. You can pass only the function name. If you need additional parameter to the function (e.g. mean and sd with dnorm), specify x where you pass the independent variable.

Results not shown: try them!

```
curve(dnorm) # Plots with x between 0 and 1
curve(dnorm, from = - 1, to = 1) # set x axis limit
curve(dnorm(x, mean = 0.2, sd = 0.5)) # Set mean and sd
curve(dnorm, add = TRUE) # works like lines/points
```

## **Some tricks to speed up coding**

---

# Keyboard shortcuts

- ALT+SHIFT+K show all keyboard shortcuts
- CTRL(CMD)+SHIFT+N open new script
- CTRL(CMD)+W close script
- ALT+- insert "<-"
- CTRL+1 move cursor to script
- CTRL+2 move cursor to console
- CTRL(CMD)+S save script
- CTRL(CMD)+SHIFT+C comment line/selected lines
- ALT+UP/DOWN move line upwards/downwards
- CTRL(CMD)+F find/replace

## Other tricks

- CTRL(CMD)+F with lines selected: if you do replace all, you replace only in the selected lines
- Try to hold ALT, drag the cursor vertically in a script and type/delete with the keyboard! Useful with aligned lines of code!
- Double click on a word/variable to select it, triple click to select whole line, four clicks (OR CTRL+A) to select all
- Use CTRL(ALT)+left/right to move rapidly word by word
- Use auto-completion!!! if you want to load a file in the working directory, type a quote (" or ') and type TAB



## References

---

- James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2013. *An Introduction to Statistical Learning*. Vol. 112. Springer.
- Wickham, Hadley. 2014. *Advanced R*. Chapman; Hall/CRC.